

Refactoring the GRC pipeline

What changed in RP7, and why

Emilia Tjernström

2026-05-04

The pipeline had drifted into twenty-two .do files we could not audit

- Twenty-two numbered .do files; experience and birth families spread across six of them
- A second Dropbox folder, `ReplicationPackage6 - real values/`, kept manually in sync with the nominal-values one
- Ster filenames colliding across .do files, with several too long for Stata's internal name limit

None of this changes the published numbers. All of it makes the pipeline harder to maintain.

We collapsed twenty-two .do files into thirteen

Before (RP6, 22 files)

- 5_GrRC.do
- 6_GrRC_NonAg.do
- 8_GrRC_hukou.do
- 10_GrRC_experience.do
- 11_GrRC_max_experience.do
- 12_GrRC_experience_share.do
- 13_GrRC_max_experience_share.do
- 14_GrRC_NonAg_experience.do
- 15_GrRC_birth.do
- 16_heterogeneity_tables.do
- ...

After (RP7, 13 files)

- 5_GrRC.do
- 6_GrRC_NonAg.do
- 8_GrRC_hukou.do
- GRC_extras.do (experience family dispatcher)
- make_tables.do, make_figures.do
- 16_heterogeneity_tables.do

Net: —9 files, same numbers

Experience and birth families: one program, one dispatcher

The five experience .do files plus `15_GrRC_birth.do` all did the same thing: GRC regression with one extra regressor.

What changed

- Shared regression logic: `run_grc_with_extra_regressor` in `0_programs.do`
- Per-cell calls: `GRC_extras.do`, one cell per (country, spec3, regressor)
- Tables: `extras_tex_table` reads sters and emits the family .tex files

To add a new family member

- One cell in `GRC_extras.do`
- One arm in `run_grc_with_extra_regressor`

One table-builder for all GRC tables

We had multiple table-building programs that were 95% the same code:

- `grc_tex_table_trend`
- `grc_tex_table_trend_exp`
- `grc_tex_table_trend_birth`
- `grc_tex_table_trend_hukou`

Now: one `grc_tex_table_trend`, four call signatures.

To change how a GRC table renders—column widths, indicator-row order, postfoot string, decimal precision—edit `grc_tex_table_trend` in `0_programs.do`.

One switch handles nominal and real values

Before

- Two parallel Dropbox folders
- Every edit propagated by hand

Now

- One global, `values(nominal|real)`, in `0_path_config.do`
- `nominal`: filenames as before
- `real`: appends `_r` suffix; both modes coexist in one `output/` tree

To switch modes: edit global `values` in `0_path_config.do`.

How ster filenames are built

Pattern: `grc_<country>_<spec3>_<covs2>[_<sfx1>].ster`

country	IDN, CHN, TZA, or CHN_hukou_<subgroup>
spec3	cuu, cub, iuu, cnu
covs2	c0, ct, c1, c2, ca
sfx1	n (never), a (always), d (delta), g (average)

Every ster filename is built from these tokens. To find the ster behind a particular cell, assemble the tokens.

Full convention in `RP7/scripts/STER_NAMING.md`.

Editing the GRC tables

To change caption, label, or notes

- Edit the `\GRCtable{}` and `\GRChukoutable{}` macros in `preamble.tex` (the paper)

To change which coefficients appear, label text, or formatting

- Per-cell options: `make_tables.do`
- Underlying template: `grc_tex_table_trend` in `0_programs.do`

Cosmetic refresh in this round

- Average Δ row reads $\bar{\Delta}$
- Spurious blank rows from `esttab` stripped after output

What this buys us

- One source of truth for the regression code
- One source of truth for the table caption, label, and notes (in the paper preamble)
- Ster filenames you can read, type, and look up
- Same published numbers as before

The pipeline is now small enough to hold in one head.