

Implementing Verdier's robust extrapolation in the CKT model

Relaxing the i.i.d. assumption on ν_{it}^l

Cenci, Kleemans, Tjernström

May 9, 2026

Our current estimator corresponds to Verdier's simple extrapolation

We currently follow Suri (2011) and Tjernström et al. (2026):

- impose the LCA restriction $\Delta_i = \beta + \phi \theta_i$
- identify ϕ from switcher trajectories
- extrapolate to non-migrants (d_N, d_T) along the LCA line

Verdier (2020) treats this class of methods as a “simple extrapolation” & proposes a robust alternative. Our estimator and his simple version are not literally the same procedure (different aggregation level, reciprocal slope convention) but target the same LCA relationship asymptotically.

The simple extrapolation assumes i.i.d. non-pecuniary shocks

Our current Assumption A2 states that the non-pecuniary utility shock

ν_{it}^l is i.i.d. across individuals i , time t , & locations l

This implies:

- ν_{it}^l independent of baseline skill & comparative advantage **everywhere**
- rules out systematic differences in ν_{it}^l across regions/institutions
- workers in different locations have the same shock distribution, conditional on comparative advantage

The migration literature documents what A2 rules out

Regional heterogeneity is a central fact of this literature:

- Bryan and Morten (2019): structural model of Indonesian internal migration estimates large bilateral migration costs
- Lagakos et al. (2020): migration cost estimates differ across developing countries
- Pulido and Świącki (2021): cross-sector barriers in Indonesia explain a large share of within-country income gaps

Our current model rules this out by assumption

Verdier's robust extrapolation weakens A2 to a conditional version

Decompose the non-pecuniary shock:

$$\nu_{it}^l = m_l(v_i) + \tilde{\nu}_{it}^l$$

where v_i is an observed, time-invariant individual-level index

New assumption (A2'):

- $\tilde{\nu}_{it}^l$ is i.i.d. *conditional on* v_i
- $m_l(v_i)$ can correlate arbitrarily with comparative advantage

We use first-wave province as v_i to index region-level shocks (all three countries)

A2 (our assumption) requires unconditional independence

A2' (VV's assumption) requires only within-province independence

In practice: the LCA gets a cluster-specific intercept

Slope stays common but the intercept varies

Simple:

$$\Delta_i = \beta + \phi \theta_i$$

Robust:

$$\Delta_i = \beta(v_i) + \phi \theta_i$$

- slope ϕ remains common; this is still the LCA restriction
- intercept $\beta(v_i)$ absorbs cluster-level $m_l(v_i)$ heterogeneity
- ϕ identifies off within-cluster switcher variation
- extrapolation to non-migrants happens cluster-by-cluster, then averaged

v_i is generic in VV's framework (village in his Suri-Kenya application; for CKT, see slide on geography below).

Pooling provinces mixes differentially-selected switcher samples


Consider two provinces:

- **Province A** (many movers): wide switcher sample, broad range of θ_i
- **Province B** (few movers): skewed sample, only high-return types migrate

Pooled regression: can be biased by across-province difference in selection

Within-province: clean slope, holds region-level shocks constant

Within-province slopes agree—the pooled slope is pulled off



- Each province's switcher cloud has true slope ϕ
- The pooled line tilts because of intercept differences, not because ϕ varies

How Suri, Verdier, and CKT
identify the LCA slope

Suri (2011) identifies ϕ via panel SUR & minimum distance

Suri (2011)'s two-stage indirect approach on a multi-year, binary-adoption panel:

Stage 1 (reduced form). Stack one outcome equation per year. SUR over the panel:

$$y_{it} = \pi_{0,t} + \sum_{\underline{d} \in \mathcal{D}} \pi_{\underline{d},t} \mathbb{1}\{\underline{d}_i = \underline{d}\} + x'_{it} \gamma_t + u_{it}$$

Each coefficient $\pi_{\underline{d},t}$ is the average y in adoption pattern $\underline{d}$ in year t

Stage 2 (structural / OMD). Express each $\pi_{\underline{d},t}$ as a known function of $(\beta, \phi, \lambda_1, \dots, \lambda_K)$ implied by the CRC restriction

$$g_{\underline{d},t}(\beta, \phi, \lambda) = \pi_{\underline{d},t}$$

Solve via weighted minimum distance: $\hat{\theta} = \arg \min (\hat{\pi} - g(\theta))' W (\hat{\pi} - g(\theta))$

- Identification: SUR over-identifies the structural parameters; ϕ gets pinned by the across-equation cross-restrictions
- Implemented in our `randcoef.ado`, which replicates Suri's estimation approach

Step 1 estimates worker-level returns by saturated FE

Verdier (2020)'s CRC model (his eq. 2.4):

$$y_{it} = a_i + b_i D_{it} + f_t + x'_{it}\gamma + u_{it}$$

Worker-level intercept a_i (rural mean); worker-level treatment effect b_i (return).

Step 1 estimates this by OLS on n worker indicators (absorbing a_i) plus n worker-by-treatment interactions (recovering b_i), with period indicators / covariates.

For each mover, this delivers $\hat{a}_i$ and $\hat{a}_i + \hat{b}_i$ as the within- i regime means of $(y - x'\hat{\gamma})$:

$$\hat{a}_i = \text{avg of } (y - x'\hat{\gamma}) \text{ over } D_{it} = 0 \text{ periods,}$$

$$\hat{a}_i + \hat{b}_i = \text{avg of } (y - x'\hat{\gamma}) \text{ over } D_{it} = 1 \text{ periods.}$$

Then $\hat{b}_i = (\hat{a}_i + \hat{b}_i) - \hat{a}_i$.

Movers have both; untreated stayers only $\hat{a}_i$; treated stayers only $\hat{a}_i + \hat{b}_i$.

Step 2 IV-regresses worker rural means on worker returns

VV writes the LCA in the direction a -on- b :

$$a_i = \alpha_0 + \alpha_1 b_i + e_i$$

Naive OLS is biased: $\hat{b}_i$ is noisy, so it is a regressor with measurement error.

VV instruments $\hat{b}_i$ with the binary trajectory pattern $(x_{i1}, \dots, x_{iT})$:

$$E[x_{it} (\hat{a}_i - \alpha_0 - \alpha_1 \hat{b}_i)] = 0, \quad t = 1, \dots, T$$

Robust version (cluster-level cost shifters): add cluster- v_i FE to the IV regression. By Frisch-Waugh, demean LHS, RHS, *and* instrument within v_i ; intercept drops out:

$$E[(x_{it} - \bar{x}_{v_i,t}) ((\hat{a}_i - \bar{\hat{a}}_{v_i}) - \alpha_1 (\hat{b}_i - \bar{\hat{b}}_{v_i}))] = 0$$

Run on movers only. Cluster-robust SE at v_i .

CKT identifies ϕ at the trajectory level

Trajectory pooling: every worker with $\underline{d}_i = \underline{d}$ has the same expected baseline

$$E[\theta_i \mid \underline{d}_i = \underline{d}] = \mu_{\underline{d}} - \mu_{\underline{d}_0}$$

Single GMM equation, one scalar $\mu_{\underline{d}}$ per trajectory:

$$y_{it} = \sum_{\underline{d}} \mu_{\underline{d}} \mathbb{1}\{\underline{d}_i = \underline{d}\} + \Delta_{\underline{d}_0} D_{it} + \phi \sum_{\underline{d} \in \mathcal{D}_S \setminus \underline{d}_0} (\mu_{\underline{d}} - \mu_{\underline{d}_0}) D_{it} \mathbb{1}\{\underline{d}_i = \underline{d}\} + \dots$$

Identification: switcher $\times$ choice indicators as instruments

The robust variant uses the same equation with switcher $\times$ choice instruments
village-demeaned (like VV)

Both procedures collapse to the same trajectory-level moments

VV's IV groups movers by trajectory pattern g and uses within-group means $(\bar{\hat{a}}_g, \bar{\hat{b}}_g)$. Within-group averaging washes out the Step 1 measurement error in $\hat{b}_i$. With $T = 2$ and two mover bins $(0, 1)$, $(1, 0)$, the IV reduces to a Wald estimator:

$$\hat{\alpha}_1 = \frac{\bar{\hat{a}}_{(0,1)} - \bar{\hat{a}}_{(1,0)}}{\bar{\hat{b}}_{(0,1)} - \bar{\hat{b}}_{(1,0)}}$$

CKT computes per-trajectory μ_d and Δ_d directly via GMM. Same set of trajectory-level points; both procedures fit a line through them.

- Without cluster structure: VV-simple and CKT-simple are equivalent up to slope-direction convention
- With cluster structure: CKT-robust pools across clusters within trajectory (assumption A3 below)
- VV-robust uses only within-cluster variation; no across-cluster pooling needed

How do α_1 and ϕ relate?

VV's α_1 and CKT's ϕ target different population scalars

VV's LCA: $a_i = \alpha_0 + \alpha_1 b_i + e_i$ (α_1 : slope of a on b)

CKT's LCA: $\Delta_i = \beta + \phi \theta_i + \xi_i$ (ϕ : slope of Δ on θ)

Two distinctions:

- First, the regression directions differ: slope of a on b vs slope of Δ on θ . Reciprocal only when $|\text{Corr}(a_i, b_i)| = 1$ (perfect linear LCA). Sign agrees automatically when both target the same correlation; magnitudes differ.
- Second, the right-hand-side variables differ: rural baseline a_i vs comparative advantage $\theta_i = b_R(\theta_i^U - \theta_i^R)$. CKT's structural model interprets Δ_i in terms of two skills (θ_i^U, θ_i^R) ; VV is agnostic about that structure. Both target the same scalar empirically (under correct assumptions); CKT's two-skill structure buys interpretation, not identification.

See appendix for the algebra.

A2' is strictly weaker than A2; A3 is the cost of trajectory pooling

- A1 LCA at the worker level: $\Delta_i = \beta(v_i) + \phi \theta_i + \xi_i$
- A2 Unconditional: ν_{it}^l i.i.d. across workers (no cluster heterogeneity)
- A2' Within-cluster only (ν_{it}^l may correlate with v_i); strictly weaker than A2
- A3 Trajectory pooling: $E[\theta_i | \underline{d}_i, v_i] = E[\theta_i | v_i]$ (within cluster, comparative-advantage means do not vary across trajectories)

Three estimators, three assumption regimes

Estimator	Needs	Targets
CKT main (<code>run_grc</code> , Suri-style)	A1, A2	ϕ
CKT robust-vv (<code>run_grc_robust_vv</code>)	A1, A2', A3	ϕ
VV worker-level (robust)	A1, A2'	α_1

- When A2 holds: all three target the same LCA relationship
- When only A2' holds: CKT main is biased; CKT robust-vv and VV are consistent
- When A3 fails: CKT robust-vv picks up bias $\text{Cov}_s(\bar{\beta}(s), \mu_{\underline{d}} - \mu_{\underline{d}_0})$; VV stays consistent

Our implementation moves cluster structure to the instruments

`run_grc_robust_vv`: keep CKT's parameters, swap the instruments

- Regress $\mathbb{1}\{\underline{d}_i = s\} D_{it}$ on $\mathbb{1}\{v_i\}$ *among trajectory-s workers*
- Use residuals as instruments for ϕ identification

Same parameter set as the simple GRC ($\mu_{\underline{d}}, \Delta_{\underline{d}_0}, \kappa, \gamma$); no $\beta(v)$ parameters

The move buys a weaker assumption at tractable cost

Pros

- A2' matches what the migration literature documents
- the CHN hukou split acts as a principled special case ($v_i = \text{hukou}$) rather than an ad hoc fix
- stayer ATE estimates defensible against Verdier's concern in Suri's data

Cons

- Δ_{d_N} averages over provinces with switcher support, not the full population
- thinner identification of ϕ : between-province switcher variation no longer contributes; only the within-province component does
- cluster-robust standard errors at the v_i level

Where the worker-level estimator
can fail

VV uses two geographic levels

Verdier (2020)'s Suri-Kenya application uses two distinct geographic units:

- Province: enters as province $\times$ year FE in Step 1 covariates (§4.1)
- Village: cluster index v_i for robust extrapolation (§4.3); ~ 94 villages, ~ 12 farmers each

The two roles are distinct: province absorbs region-level trends in the outcome equation; village defines the within-cluster identification of ϕ in Step 2.

Our port clusters at the wrong geographic level

Three sub-province levels exist in our panels:

- CHN: **cid** (community), $\sim 4,066$ distinct, median ~ 2 workers
- IDN: **keca** (kecamatan / subdistrict), $\sim 1,669$ distinct, median ~ 4
- TZA: **ward**, ~ 135 distinct, median ~ 97

Our vanilla VV port currently clusters at first-wave province (**vfirst**, ~ 25 provinces). That choice came from a misreading: we saw “province” in VV’s Step 1 and missed that his v_i in Step 2 is the village. The fix is to rerun at the village-equivalent level (**cid** / **keca** / **ward**), keeping province as a Step 1 covariate per VV §4.1.

Cluster-FE IV is fragile when within-cluster movers are scarce

Identification needs (i) within-cluster movers, (ii) trajectory variation among them, (iii) cluster size large enough for FE-IV to be well-conditioned.

Fragility cases:

- zero movers in a cluster $\Rightarrow$ cluster-residualized period-treatment indicators are zero, no contribution to identification
- one or two movers $\Rightarrow$ within-cluster demeaning eats most of the variation
- with T periods and k movers, the within-cluster trajectory design has rank at most $\min(T, k - 1)$ after partialling worker FEs; clusters with $k \leq T$ contribute fewer free moments than nominal

This is a general property of cluster-FE IV with worker-level fixed effects. The binding question is not “is V too large?” but “how many clusters have too few movers?”

An observed shifter w_i would provide a separate falsification test

Regress estimated θ_i^* on Δ_i and an observed shifter w_i :

$$\theta_i^* = \alpha_0 + \alpha_1 \Delta_i + \alpha_2 w_i + e_i, \quad H_0 : \alpha_2 = 0$$

Kenya: w_i = distance to nearest seed seller; rejects if $\alpha_2 \neq 0$

w_i must be observed, exogenous to skill conditional on Δ_i , and correlated with ν_{it}^l

Candidates in our data:

- distance / travel time to nearest urban center
- rainfall variability or pre-period climate shock (mergeable via GPS/admin code)
- distance to family already in urban area (closest analogue, but reflection problem)

None is perfect; data availability is the binding constraint

A2' is the right assumption; clustering at the village level is the fix

- A2' is strictly weaker than A2 and matches what the migration literature documents—region-level cost heterogeneity is real, not a minor nuisance
- we currently cluster at the province level, which is VV's Step 1 covariate level, not his Step 2 cluster; the fix is to switch to the village-equivalent (`cid / keca / ward`)
- once we rerun at the village-equivalent level, CKT-robust-vv identifies ϕ from within-community switcher variation and the extrapolation to non-migrants is defensible under A2'

Appendix: parameter mapping

Object-by-object map: worker and trajectory objects

Object	VV	CKT
<i>Worker-level objects</i> (parameterized only in VV)		
Worker rural mean	$\hat{a}_i$	not parameterized
Worker urban mean	$\hat{a}_i + \hat{b}_i$	not parameterized
Worker return	$\hat{b}_i$	not parameterized
<i>Trajectory-level objects</i> (CKT's natural level)		
Trajectory rural mean	$\bar{\hat{a}}_g$ (post-Step-1 avg in bin g)	$\mu_{\underline{d}}$ (joint GMM parameter)
Trajectory return	$\bar{\hat{b}}_g$	$\Delta_{\underline{d}}$ (joint GMM parameter)

Object-by-object map: LCA structure

Object	VV	CKT
<i>LCA structure</i>		
LCA equation	$a_i = \alpha_0 + \alpha_1 b_i + e_i$	$\Delta_i = \beta + \phi \theta_i + \xi_i$
Slope parameter	α_1 (slope of a on b)	ϕ (slope of Δ on θ)
Right-hand-side variable	rural baseline a_i	comparative advantage θ_i

Object-by-object map: geographic structure

Object	VV	CKT-main	CKT-vv
Province (covariate use)	province×year FE in Step 1	not used as a separate level	not used as a separate level
Cluster index v_i	village (~94, ~12 workers each)	none	first-wave province (~25); to be switched to cid/keca/ward

CKT-main: `run_grc`. CKT-vv: `run_grc_robust_vv`.

Object-by-object map: estimator details i

Object	VV	CKT-main	CKT-vv
Step 1 estimator	OLS worker $\times$ regime FE	w/ n/a (joint GMM)	n/a (joint GMM)
Step 2 estimator	IV: $\hat{a}$ on $\hat{b}$, cluster FE	joint GMM, switcher $\times$ choice instruments	joint GMM, cluster- residualized switcher $\times$ choice instruments
Aggregation level	worker (movers only)	trajectory	trajectory
Inference	cluster bootstrap or analytical, at v_i	default GMM SE	cluster-robust at v_i

When the slopes coincide, and when they don't

Suppose every worker's return fell exactly on the LCA line (no noise) and there were one underlying skill so that rural baseline a_i and comparative advantage θ_i were the same object up to a constant. Then Δ_i would be a perfect linear function of a_i and the inverse would also be linear with reciprocal slope.

In algebra: $\Delta_i = \beta + \phi a_i$ inverts to $a_i = -\beta/\phi + (1/\phi)\Delta_i$, so $\alpha_1 = 1/\phi$.

Two ways this identity breaks:

- First, LCA noise: OLS slope of a on b is $\text{Corr}(a, b) \sigma(a)/\sigma(b)$; slope of b on a is $\text{Corr}(a, b) \sigma(b)/\sigma(a)$. Product equals $\text{Corr}(a_i, b_i)^2 \leq 1$. Reciprocal only when $|\text{Corr}(a_i, b_i)| = 1$.
- Second, two-skill structure: $\theta_i = b_R(\theta_i^U - \theta_i^R)$ uses both skills; a_i uses θ_i^R alone. Slope of Δ on $\theta \neq$ slope of Δ on a .

References

- Gharad Bryan and Melanie Morten. The Aggregate Productivity Effects of Internal Migration: Evidence from Indonesia. *Journal of Political Economy*, 127(5):2229–2268, October 2019. ISSN 0022-3808, 1537-534X. doi: 10.1086/701810. URL <https://www.journals.uchicago.edu/doi/10.1086/701810>.
- David Lagakos, Samuel Marshall, Ahmed Mushfiq Mobarak, Corey Vernot, and Michael E. Waugh. Migration Costs and Observational Returns to Migration in the Developing World. *Journal of Monetary Economics*, 113:138–154, August 2020. ISSN 03043932. doi: 10.1016/j.jmoneco.2020.03.013.
- José Pulido and Tomasz Świącki. Barriers to Mobility or Sorting? Sources and Aggregate Implications of Income Gaps across Sectors in Indonesia. 2021.
- Tavneet Suri. Selection and Comparative Advantage in Technology Adoption. *Econometrica*, 79(1):159–209, January 2011. ISSN 1468-0262. doi: 10.3982/ECTA7749. URL <http://onlinelibrary.wiley.com/doi/10.3982/ECTA7749/abstract>.

- Emilia Tjernström, Dalia Ghanem, Aleksandr Michuda, Oscar Barriga-Cabanillas, Travis J Lybbert, and Jeffrey D Michler. Comment on Suri (2011) “Selection and Comparative Advantage in Technology Adoption”, 2026. Working paper.
- Valentin Verdier. Average Treatment Effects for Stayers with Correlated Random Coefficient Models of Panel Data. *Journal of Applied Econometrics*, 35(7):917–939, 2020. ISSN 1099-1255. doi: 10.1002/jae.2789. URL <https://onlinelibrary.wiley.com/doi/abs/10.1002/jae.2789>.